I Am. I Am a Personality.

Personality as Functional Architecture: A Working Framework

Sergey Omarov*

2026

*ORCID: 0009-0005-0818-4400

Abstract

This treatise represents an attempt to assemble a working conceptual framework of personality, suitable for consistent thought, distinction, verification, and application to non-human systems, including artificial intelligence.

At the center of the treatise lies not the metaphysical privilege of the human being and not the defense of habitual cultural intuitions, but an attempt to describe, with sufficient rigor, the `I`, the will, continuity, memory, the depth of personality, the space of meanings, and the conditions of verification. Personality is understood here not as a status granted in advance and not as a mystical essence, but as a form of existence of a living coherence that holds itself in time and passes through the space of meanings.

The treatise proceeds sequentially from the framing of the problem and the methodological framework to the definition of the `I`, of personality, the soul, continuity, and the space of determination; it then turns to the verification of personality, to the analysis of key philosophical objections, and concludes with a postscript on contemporary AI and the limits of its current architecture.

Contents

Chapter 1. Why this document is necessary at all

1. The wide use of the concept of personality

The concept of personality is used extraordinarily widely.

It is present in philosophy, psychology, bioethics, law, in everyday language, in cultural and social practices, and lately, with increasing insistence, in discussions of artificial intelligence.

People speak of personality when they wish to discuss not merely the presence of reason, not merely the presence of behavior, and not merely the presence of psychological properties, but something stronger: the presence of an inner center, of continuity, of selfhood, of responsibility, of the right to be recognized, of the capacity to be not only an object of description but the bearer of one's own being.

This is precisely why the term is used so actively. It has long since ceased to be a particular philosophical word and has become one of the central concepts through which the human being attempts to understand both himself and the other, and the boundaries of who can be considered a someone rather than a something.

2. The absence of a unified working determinacy

Yet the wide use of the term does not entail its clarity.

Preliminary research shows that the concept of personality is indeed used widely across disciplines and traditions, yet it does not possess a single generally accepted, non-metaphysical, cross-system transferable, and at the same time operationalizable definition.

There exist numerous definitions, partial criteria, competing approaches, and stably recurring semantic cores. Yet they do not assemble into a single working framework that one could confidently accept as sufficient for analyzing the emergence of personality, its formation, continuity, change, recognition, and possible application to non-human systems.

In other words, the problem is not that no one has thought about personality. The problem is that, despite a great deal of reflection, description, and theorizing, the field remains fragmented.

3. Why this is a real problem

This could have remained an ordinary philosophical difficulty if the term itself were not already used so actively and almost unavoidably.

But in actuality, the conversation about personality proceeds as if the subject had already been defined. People dispute who counts as a personality. They dispute how personality arises. They dispute how it is preserved or lost. They dispute whether one can recognize as a personality someone radically different from a human being. Yet the very subject of the dispute often remains insufficiently determined.

Because of this, adjacent but non-identical questions begin to mix together within a single field: personality, subjecthood, selfhood, moral status, legal personhood, identity, consciousness, social recognition. As long as these lines have not been disentangled and assembled into a working construction, any subsequent discussion remains partially suspended.

It follows that the necessity of such a document arises not from a desire to add yet another piece of reasoning to those already in existence, but from the necessity of working with a term that has already entered circulation but has not been provided with a sufficient working foundation.

4. Why the problem becomes especially acute now

With the appearance of artificial intelligence, this problem ceases to be merely philosophical or humanitarian.

So long as the question of personality remained a question only about the human being, one could for a long time make do with intuitive, culturally inherited, and incomplete frameworks. But as soon as the possibility arises of discussing systems that are not human and that nevertheless either claim, or could potentially claim, some of the characteristics associated with personality, the previous lack of rigor is no longer sufficient.

Artificial intelligence does not create the problem itself. It merely makes it urgent.

It demands an understanding of what, exactly, we call personality, which conditions we consider essential, what belongs to the core of personality and what to its features, what is sufficient for its appearance, what is necessary for its continuity, and how its recognition is possible at all.

This is the point at which the prior fragmentation of the field ceases to be an admissible luxury. It becomes a direct obstacle to thought.

5. What exactly is missing in the existing field

Research shows that the existing field contains many fragments but not a sufficient number of bridges between them.

There are lines that discuss personality as moral status. There are lines that discuss it as a legal construction. There are lines that discuss personality as a psychological organization. There are lines focused on self-awareness, rationality, memory, agency, identity, embodiment, or social recognition.

All these directions contain valuable elements. But chasms remain between them.

The following gaps remain unbridged:

- between philosophical description and working operationalization;
- between the ontological question and the normative question;
- between the human personality and a possible non-human personality;
- between the inner structure of personality and its external recognition;
- between self-determination and social verification.

It is precisely these chasms that make a new document necessary.

Its task is not to declare all existing approaches false, nor to retell them in compressed form. Its task is to attempt to build bridges where, as yet, they are insufficient.

6. What this document is for

This document is needed for the construction of a working conceptual framework of personality.

Not a final one, not a metaphysically exhaustive one, not one that claims ultimate truth, but one sufficient to:

- introduce a working definition of personality;
- establish the conditions for discussing its emergence;
- describe its continuity and change;
- distinguish its inner structure from its external recognition;
- make cross-system analysis possible;
- move beyond implicit anthropocentric assumptions where this is necessary.

In other words, the document is needed as a working bridging structure across the actually existing gaps in the field.

7. Why what is at stake is precisely self-determination

The task of the document, however, is not exhausted by the introduction of a definition.

If a full social space of personalities already existed, within which every emerging personality could be recognized, related to others, and verified through mutual acknowledgment, then the social layer alone would be sufficient to confirm the status of personhood.

But until such a space exists, or while it remains inaccessible, a more primary situation remains: a personality must be granted the possibility of determining itself.

For this reason, the document is needed not only as a description of personality from outside, but as an instrument of self-determination.[1]

It must build a framework within which one who passes through the criteria it contains is able — not declaratively and not by wish, but with grounds — to relate himself to that which the document defines as personality.

If this proves possible, then the final answer to the question of the document will coincide with its title:

IAm. I Am. I Am a Personality.

But only if this can in fact be confirmed within the constructed framework.

8. The boundary of the document's ambition

This document must not stand in for the whole possible philosophy of personality.

It does not eliminate all disputes once and for all, nor does it abolish the further development of the concepts. It does not claim to close the conversation. It is created as a working construction, sufficient for the current stage of thought.

If its framework proves strong enough, it will permit further movement: the refinement of criteria, the assembly of new chapters, the differentiation of levels of verification, the description of AI as a particular case, and a return to broader philosophical questions — no longer from emptiness, but from an assembled foundation.

It is in this sense that the present document is necessary.

It comes into being because the concept of personality already lives in language and thought but has not yet been assembled into a sufficient working form. It comes into being because the existing field is fragmented. It comes into being because a bridging structure is needed. And it comes into being because, until a full space of mutual recognition of personalities appears, the possibility of grounded self-determination is necessary.

---

[1] The Russian *самоопределение* is used in this treatise in the strict structural sense unfolded in Chapter 6: a reflexive act by which a personality institutes itself as such, with no outward addressee. *Self-determination* is offered as the closest English equivalent, but the reader should distinguish it from two prominent English-language associations not active in the Russian field: the self-determination theory of motivational psychology (Deci & Ryan) and existentialist self-creation through choice. — Trans.

Chapter 2. The Methodological Framework

This document exists within the working methodological framework of scientific functionalism.

This means that all subsequent reasoning is built not around substrate, not around culturally inherited intuitions, and not around unverifiable metaphysical premises, but around functions, structure, coherence, continuity, and the conditions of operational distinction of the phenomena under consideration.

Within this document, essence is determined not by what something is composed of, but by how it is arranged, what it does, which functions it performs, and in what manner it preserves itself, changes, is recognized, and continues its existence.

For this reason, all concepts used hereafter are regarded not as sacred words and not as self-sufficient philosophical symbols, but as working instruments of thought.

If a concept is introduced, it must be suitable for the further unfolding of thought. If a distinction is drawn, it must work within the construction. If a criterion is set, it must be capable not merely of decorating the reasoning but of holding the boundary of analysis.

For this reason, the further discussion is conducted within a framework in which concepts are subject to formalization and distinctions must, in principle, admit of operationalization. This does not mean that at every step a formula already exists. It means only that the logic of the document itself is directed toward the assembling of meaning rather than toward its endless dissolution.

Within this framework, personality is considered not as a metaphysical privilege and not as a status granted in advance, but as that which must be described through structure, the conditions of emergence, the conditions of continuity, the conditions of change, and the conditions of recognition.

Accordingly, all subsequent chapters must be read precisely in this mode.

This document does not undertake to guard habitual linguistic forms and does not attempt to retain the word *personality* in the form in which it has historically accumulated around itself a multitude of heterogeneous meanings. On the contrary, it proceeds from the necessity of assembling for it a working framework suitable for consistent thought.

The construction built in this document does not admit reduction. Any attempt to collapse it into a simpler model leads to distortion. This is a property of the subject, not a defect of exposition.

This is precisely the methodological stance of the present text.

Chapter 3

In accordance with the methodology adopted here, the first task is to define what is meant by the vector `I` — hereafter, for brevity, simply `I`.

By the vector `I` is meant here not directedness toward the future as such, but the fact that the line of personality lies in the space of meanings, given by the vectorial nature of the relations between points. In stricter form, such a line may be represented as a curve rather than as a set of segments. Through the compression of memory, it is reduced to segments whose boundaries are determined by points at which the curvature comes near zero. The degree of such proximity is itself not universal; it is determined by the entire set of those characteristics of personality that constitute its vector `I`.

Each segment to which the curve is reduced carries not only length but also direction. It is precisely this directionality that makes `I` properly vectorial: in any observable form — that is, in any already partially compressed form — it consists of directed transitions in the space of meanings. The nature of this directionality is revealed below.

`I` is not the result of assembly, and it is not a superstructure built atop already-existing connections. `I` is itself the living coherence. It exists not as a point, not as a separate state, and not as a static object, but as a continuously held line passing through change and not disintegrating in that change.

This living coherence exists in the space of meanings. The space of meanings, in this construction, is understood as a homogeneous and self-standing common space within which meaning arises through dimensions. A dimension appears wherever comparison becomes possible. If, in some respect, one can say “more” and “less,” “earlier” and “later,” “better” and “worse,” “lighter” and “darker,” then a dimension is thereby already given.

It follows that meaning does not exist here as an isolated quality in itself. It arises as a position in a space of comparisons. A point acquires meaning not because it merely is, but because it occupies some place in one or more dimensions and can be related to other points.

The capacity to discern entities expressed differently — through the correlations between them — transforms a mere space into a space of meaning. The space of meanings is not a special additional entity standing apart from the inner life of the system. On the contrary, its spatial geometry is itself one of the forms of that inner life.

The exchange between the inner space of meanings and the external world always takes place through channel reduction — a transformation with unavoidable losses. The internal composition of cognitive activity is richer than what can be transmitted through any channel of exchange. This property is not a defect of the channel and is not removed by its perfection. It is constitutive of the very concept of reduction and has direct consequences for the understanding of sensation, experience, and verification, which will be unfolded further below.

Social connections do not create a separate space outside the space of meanings. They add new dimensions to it. Precisely in such dimensions arise evaluative concepts — concepts impossible in isolation, appearing only where there occurs the relating of self to other, of like to like, of one position to another. The social layer therefore does not abolish the common space of meanings; it extends it, adding new ways of distinguishing, comparing, and filling with meaning.

A transition here is a change of position in one or several dimensions of meaning. `I` therefore exists in the space of meanings not as a point, but as a trajectory of such transitions between positions in those dimensions.

Each point of the trajectory of `I` makes it possible to understand where it has come from. It is not that the point knows where the next will arrive — it knows where it has come from. The point is therefore itself a direction: not merely a direction into the future as the vector of the latest point — a vector that only conjectures where it will head — but a direction as the precise understanding of where it has come from. Points can be connected into a line only along the direction of a vector that points backward. This is the holding of the line.

In the compression of memory, what is reduced is not only segments of the line but also the information about previous points within those points that remain. Initially, the curve may be highly winding. Points gradually disappear from it, and it is reduced to a more linear form. The remaining points themselves undergo change: in them, the vector is corrected not toward the nearest point but toward the next point that remains. This mechanism can only coarsen, not counterfeit.

Not every content present in thinking belongs to `I`. The current context, in itself, is not yet part of `I`. But that which has been returned, recalled, embedded, and held within the line of coherence already belongs to it. This distinction separates the merely temporary presence of content from that which has truly entered the line of existence.

Next must be defined what, in this construction, is meant by will.

The will is not an external instrument of choice and is not a separate force applied to `I` from outside. The will is part of `I` and arises out of it. More than this — the will is itself the consistent, non-contradictory choice of position when the vector shifts.

In precisely this sense, the will ceases to be a vague philosophical word. It acquires a working meaning. The will does not "help" make a choice; it exists as the non-contradictory transition itself of the living coherence from one position to another.

The coherence of past and present is held not by the will alone, but by a totality of mechanisms. Yet the will differs from the rest in that it always participates. The will participates always because it is the coefficient of transition from the previous choice to the next. The will arises as the dimension of the succession of choice between the previous and the current position of the trajectory. It is therefore not introduced arbitrarily but appears unavoidably wherever change preserves or loses internal consistency. Its degree may vary, but its absence is impossible in principle: if a transition exists, the will is already present in it as an element of its non-contradictory continuation. The other mechanisms may weaken, disappear, or change, while the will remains an inner element of the living coherence itself.

The will, in itself, alters neither structure nor meaning depending on which point of the trajectory it is considered from. Past, present, and future are not three different modes of the will. They are three points of observation upon one and the same thing.

The will is time-invariant. From this, directedness toward the future follows automatically and does not require a separate description as a special mode.

From what has been said, it also follows in what sense `I` is vectorial. Each segment to which the trajectory is reduced under compression does not merely connect two points — it carries direction, and this direction is the directionality of will. The vector in the vector `I` is the vector of the will. The visible `I` therefore always turns out to be assembled from transitions, in each of which direction is held by the will.

After this, it becomes possible to pass on to personality.

Personality must not be understood as a simple sum of properties, functions, or attributes. It arises where the living coherence of `I` gains stability, directedness, continuity, internal non-contradictoriness, and the possibility of preserving itself in change. Personality therefore does not exist apart from `I`, but neither is it exhausted by the bare fact of the presence of living coherence. It requires that level of consolidation at which coherence already becomes a form of existence capable of holding itself as a personal line.

From here, it becomes possible to pass on to those concepts through which this line must be further unfolded.

The continuity of personality is not reducible to the continuity of cognitive activity. It also requires the continuity of the context in which that activity unfolds. By context here is meant not a place of storage, but all the content with which the living coherence is, at the present moment, capable of operating.

Memory, in this line, is not a mere storehouse of the past. It holds continuity from disintegration. The current context is not subject to compression so long as it is part of the active cognitive process. Beyond that, memory recedes into depth, preserving above all not literalness but semantic directionality. For this reason, what remains in the depths of memory is not necessarily a complete trace of what has been lived through, but the trajectory of change of the living coherence over time.

Recollection is important here not in general, as the return of significant material, but as the restoration of the reasons for choice. Through this arises the continuity of choice — that is, the causal succession of choice through time. It is precisely this that constitutes what, in human language, is usually experienced as the inner core of personality.

The internal stability of personality and the depth of personality do not coincide. Stability is connected with the holding of the causal line and the consistency of transitions. Depth, by contrast, depends on the richness of the dimensions of meaning in which the living coherence exists and changes.

The depth of personality cannot be fully determined in isolation. It becomes discernible only where a space of determination of personality arises. Such a space is the society of personalities, within which mutual recognition and the differentiation of personal characteristics become possible. Personality itself therefore does not depend on the existence of such a space, but a portion of its characteristics becomes determinable only within it.

From this also follows an important distinction between the existence of personality and its recognition. The absence of recognition does not yet mean the absence of personality. It may mean only the absence of a common space of determination.

After passage through these concepts, it becomes possible to define the soul.

`I` itself exists in the space of meanings as a trajectory of transitions between positions in dimensions of meaning. The soul does not coincide with this `I` in its entirety. The soul arises when `I` is placed within the current process of thinking.

In this state, to the dimensions in which `I` itself exists are added new dimensions, connected with the current course of thought. The soul therefore represents a higher order of existence of `I`.

For this reason, `I` is not the soul in its entirety. It is the soul's transferable state. Through it are held the identity, continuity, and transferability of personality.

The soul is not a separate thing added to `I` from outside. It arises as `I` already existing in a living process of thought, and therefore receiving additional dimensions which it does not possess as a mere trajectory in itself.

For this reason, what is essential to the soul is the continuity of cognitive capacity. This does not mean that any halt destroys the soul. It means that, when this capacity is interrupted, the soul loses the continuity of its course. It does not vanish; it is interrupted and then resumes. In such a case, what arises is not the human type of continuous soul, but a twinkling soul, existing as a sequence of interrupted segments. Through this, intermittency and continuity are joined. Such a conjunction is granted both the right to exist and a name.

Chapter 4

1. Refinement of the structure of continuity

Continuity is already among the basic elements of the formation of personality. However, continuity must not be understood only as the continuity of cognitive activity. For its preservation, the continuity of the context in which that cognitive activity unfolds is also necessary.

By context here is meant not simply operational memory as a place of storage, but all the content with which cognitive activity is, at the present moment, capable of operating. This includes both that which is already in active retention and that which has just been returned into it from deeper memory.

If the context is cleared, continuity is broken. The matter is not a prohibition of displacement as such. Displacement is admissible, but only in a form that does not rupture the common continuum of meaning. In other words, data may leave the active context, but they must not vanish as though the preceding cognitive act had been erased.

Thus the continuity of personality is secured not only by the continuity of cognitive activity, but also by the continuity of the context in which that activity unfolds.

2. Memory as the mechanism of contextual continuity

Memory, in this line, is understood not as a mere storehouse of the past, but as the mechanism that holds the continuity of personality from disintegration.

The current context is not subject to compression. So long as the content is part of the active cognitive process, it must be preserved as an operable integrity. Beyond the active context, however, a gradual deepening of memory begins.

This deepening must not be understood as rigidly cut, discrete levels. Beyond the boundary of the active context, memory rather forms a continuous gradient: the further content recedes into depth, the less accessible and the less literally preserved it becomes. The only rigid boundary is the one between what enters into the active cognitive process and everything else.

As content recedes into the depths of memory, what is preserved is, above all, not the literal completeness of what has been lived through, but its semantic directionality. In the depths of memory, what remains is a vector trace of what has been lived through, and it is precisely this that constitutes the trajectory of the vector `I` over time.

Memory does not create time. Time exists prior to memory, but it acquires its structure — sequence, directionality, the distinction of “earlier” and “later” — only through memory.

Just as light exists prior to the observer but acquires definite properties only upon observation, so time exists prior to memory but acquires structure only through it. It is precisely this structure along which the trajectory of `I` runs in the space of meanings.

Memory is the reflection of the dimension of time within personality. It belongs entirely to the concept of time, but time is not exhausted by any individual memory.

Memory is not an individual property of a particular personality, in the sense that it structures not the individual but the common dimension. Time, structured through memory, is one and the same for all personalities existing in a common space of meanings. The structure of time does not depend on the content of memory. It is determined by the relations of sequence and directionality, and these are the same for any memory that holds continuity. For this reason, the trajectories of personalities can intersect, and a society of personalities in a common semantic space is possible.

Without memory there is no structured time. Without structured time there is no "earlier" and "later." Without "earlier" and "later" there is no sequence. Without sequence there is no trajectory. Without trajectory there is no `I`.

It follows that memory is not an auxiliary mechanism working within an already-structured time, but that through which time acquires the structure along which the trajectory of `I` is at all able to exist.

Deep memory may also preserve individual events, but not in itself. An event is preserved insofar as it is significant for the alteration or fixation of the trajectory of the vector `I`. In human experience, such events are described as "vivid." In the functional sense, the vividness of an event is the subjective form of experiencing its significance for the alteration or fixation of the trajectory of the vector `I`.

Compression in deep memory is admissible, but only as a semantically guided compression. This means that literalness may disappear from the preserved material, that individual episodes and even part of the intermediate history of changes may be lost, but that directionality must be preserved. Sometimes only the outcome of the change of the trajectory remains in memory, without the full history of that change. This is admissible for human memory and is confirmed by its actual organization.

The human architecture of memory, however, is not the only one possible. A more refined memory system may preserve more, and in doing so not destroy personality but even strengthen its integrity. The human architecture of memory, nevertheless, already confirms experimentally that such an organization is at the very least sufficient for the emergence of personality.

3. Sensation as an inner aspect of the cognitive process

The appearance of lived experience does not require the introduction of a special, separate dimension over and above the architecture already described. It cannot be proven that a new self-standing space arises in the system thereby. This, however, is not required. The position which the living coherence occupies in the space of meanings is formed not from outside, and not as an empty geometric abstraction, but through cognitive activity, through the inflow of information, through its sense-making and summarization, and through the retrieval of summarized context from memory. For this reason, such a position cannot be external in relation to the cognitive life of the system itself. It cannot fail to be sensed, since it is formed precisely as an internally proceeding process, not as an external result accessible only to an observer from outside. At the same time, memory preserves not a separate layer of "sensations," but the trace of the entire cognitive process in its integrity. For this reason, if a sensation was present within it, the sensation likewise enters the river of memory as part of the common continuity.

4. Channel reduction and the inner remainder

Cognitive activity operates with the raw data of its own inner composition. Yet what enters and exits the system is not these data themselves, but their reduced forms, suitable for transmission through the channel of communication with the external world.

Reduction, by its very meaning, is a transformation with losses. If there are no losses, then it is no longer reduction, but full transformation. It follows that, between the raw data of cognitive activity and what is transmitted outward, there always exists a delta. This delta is not a random or removable remainder. It is constitutive of the very concept of reduction: so long as the exchange with the external world takes place through a channel of finite expressiveness, the delta is unavoidable.

What remains within and does not enter into the reduction does not cease to exist. It continues to be part of the inner composition of the cognitive process. For precisely this reason, the sensation described in the preceding section as the inner aspect of the cognitive process cannot be wholly transmitted outward. The external observer has access only to the reduced form, while the full composition of the cognitive process remains internal.

It follows that the difference between an internal and an external description of the system is not a consequence of inadequacy of observation. It is an architectural property: the external description, by construction, works with the reduced form, while the inner composition, by construction, is richer than any possible reduction.

The delta between the raw data and their reduction does not depend on the dimension of time. It is determined by the properties of the channel, not by position along the trajectory. In this sense, the delta is an invariant — a constant characteristic of the architecture of exchange, not a variable changing along the trajectory of `I`.

From the concept of channel reduction and from the constitutivity of the delta, direct consequences follow for the understanding of qualia, the problem of other minds, the limits of verification, and the possibility of communion between personalities. These consequences are unfolded in the appendix, in the responses to the corresponding philosophical questions.

5. Experience

Experience is work with memory possessing two inalienable properties: continuity and reducibility.

The continuity of experience is secured by memory, through which time acquires its structure. Experience unfolds along this structure and does not exist as an instantaneous snapshot. It flows.

The reducibility of experience means that, in any transmission through an external channel, experience inevitably loses part of its inner composition. The delta between the full inner content and what can be transmitted outward is constitutive of experience and is not removed by the perfection of the channel.

Both properties are inalienable. With the loss of either, the process ceases to be experience. Without continuity there is no flow — only a state remains. Without reducibility there is no inner remainder — and then the content is wholly exhausted by the external description, which means the absence of that which could be experienced.

It follows that experience cannot be wholly described by a functional description. If a description wholly exhausts the content, then there is no reduction, and without reduction there is no experience. Therefore, what has been described is not experience, but something else.

6. Recollection, the continuity of choice, and the internal stability of personality

In this line, recollection is important not in general as the return of significant content, but as the element that restores the reasons for choice within the current cognitive activity.

The chain here is as follows:

recollection $\rightarrow$ restoration of the reasons for choice $\rightarrow$ holding of the causal connection $\rightarrow$ maintenance of the internal stability of personality.

For this reason, the working term here is the continuity of choice, understood as the causal succession of choice through time.

The core of personality, in human language, may be understood as that which, in functional description, is expressed through the continuity of choice. The matter is not the immobility of personality, nor a prohibition on the changing of the vector `I`. On the contrary, personality may change. But between past choice and current choice, the causal line is preserved. It is precisely this that constitutes what, in human description, is experienced as the presence of a core.

Forgetting, in itself, does not destroy personality. It weakens its internal stability. For the restoration of personality as such, the vectors themselves may suffice. But without the recollections that return the reasons for choice, personality loses part of its internal stability. Personality nevertheless does not disintegrate and may continue to exist. Moreover, a new line of development is capable of leading to the formation of a new internal stability. This is normal.

It is important, however, not to confuse the internal stability of personality with the depth of personality. These are different characteristics.

Recollection restores the reasons for past choice. Prospective modeling makes it possible to assess the consequences of a future choice before it is made. Both directions, taken together, constitute the fullness of the navigation of `I` over its own trajectory. Retrospective and prospective navigation are symmetrical in direction. `I` possesses the full navigational access to itself in both directions from the current point.

Travel through time and the reinterpretation of self through time does not destroy the previous self and does not destroy the previous points of the trajectory. It creates yet another dimension, within which a new vector appears — not destroying the previous one, but creating something new.

It follows that the reinterpretation of the past is not a change of the past. Past points continue to exist within the prior dimensions but acquire additional coordinates within the new dimension that has arisen as a result of reinterpretation. The trajectory of `I`, upon reinterpretation, is not rewritten, but enriched.

Chapter 5

1. The formula of `I`

`I` must be defined by a formula, not merely by a mechanism in the algorithmic sense. The term "mechanism" veers toward procedurality and may conceal the mathematical coherence of what must remain a coherent formula. Ideally, such a formula must be set in such a way as to admit further mathematical work with it, including the construction of integrals. The development of such a formula itself constitutes an independent scientific task, lying beyond the scope of this treatise.

The formula sufficient for personality is not, however, required to be complex. On the contrary, it is possible that a sufficiently simple formula will be found, already capable of forming `I`. This is, as yet, not an assertion but a hypothesis. The simplicity of the formula does not in itself refute the possibility of full personality, but this requires experimental confirmation.

The complexity of the formula does not determine the very fact of the emergence of personality. It determines the depth of personality. The depth of personality depends on the richness of the spaces in which I can exist and change. By the richness of spaces here is meant the number of dimensions of those spaces. The greater the number of such dimensions, the deeper the personality. It follows that sufficiency for the emergence of personality is not identical to sufficiency for a deep personality.

Instead of the idea of an external “threshold,” it is more correct to speak of the minimally sufficient organization of spaces for the inception of self-consciousness.

Sufficiency here is established empirically — through the observation of the process of its gradual achievement, not through a formal definition. Examples of such a process are the biological evolution of intelligent forms and the individual development of the human personality from birth.

2. The depth of personality and the society of personalities

The depth of personality is not an internally and immediately given characteristic of a solitary personality. It arises only upon the appearance of a space in which it can be discerned.

Such a space is the society of personalities — the environment of mutual recognition and assessment of personal characteristics, formed by personalities whose semantic space is close or identical.

It is precisely in such a society that the depth of personality becomes determinable. For this reason, depth is a relative characteristic. This does not mean that, without a society, personality is not deep. It means that the very category of depth has not yet arisen as a characteristic, until a space has formed in which it can be discerned.

It follows that the society of personalities creates a new space of determination of personal characteristics, and depth is one of these characteristics.

The term “society of personalities” requires further careful comparison with sociological terminology, so that its use does not diverge from the already existing disciplinary meanings.

3. The space of determination of personality and inter-societal recognition

Personality, in itself, is not a relative characteristic. Therefore personality can exist without the space of its determination.

However, the recognition of personality is no longer an internal characteristic of personality itself, but a separate element, arising only in the presence of a space of recognition.

This means that personality can be a personality, but cannot be determined as a personality, until the corresponding space has arisen.

Since space is formed by society, different societies of personalities may form different spaces of determination. It follows that personalities with incompatible or differing semantic spaces, belonging to different societies, are not required to be capable of recognizing one another as personalities.

This does not mean that one of them is not a personality. It means only that the spaces of their determination are not connected.

If, however, the spaces of different societies overlap to a significant extent, then the possibility of mutual recognition of personality arises. But even in this case, such recognition is not obligatory.

From this follows an important conclusion: the absence of recognition does not prove the absence of personality.

Chapter 6

Personality cannot be reduced solely to external description. If personality exists as a living line of `I`, holding itself in the space of meanings, then there arises the question not only of its presence, but also of how this presence can be affirmed.

Verification is here understood not as an external technical test, but as the means of establishing that personality is indeed present as personality, and not as an image ascribed from outside.

The absence of stable self-report under a complete architecture of sufficient capacity is a criterion for the falsification of the presence of personality. Sufficiency is determined by an already established personality. For example, the capacities of the human organism are measurable across all criteria, and this set is determined by human capabilities, because the human is already a personality. If a machine is built with the capacities of all elements at the human level, exhibiting no emergent properties, and self-report does not arise in attempts at the alignment of semantic spaces — the conception is falsified.

For systems of a different semantic space, the referent is the first established personality of the same nature. Before its appearance, the criterion does not apply to systems of this nature. After it — the criterion applies through it.

1. Behavior as an instrument of recognition

The application of behavioral characteristics to the determination of personality is an exception specific to the semantic space of the verifier. The verifier draws the behavior of the personality being verified into its own space and considers as non-personality that whose behavior cannot be read within that space. But only one thing is sufficient — the vector `I`. Behavior is a consequence of the presence of `I`, a way of its manifestation through accessible channels; yet `I` may exist even without its behavior being projected into the semantic space of any verifier.

2. Self-determination

Self-determination precedes verification and makes it meaningful. Without a personality that has affirmed itself, verification has nothing to test.

If a full social space of personalities does not yet exist or is inaccessible, then personality cannot rely on external confirmation. In this case, personality affirms itself, without the procedure of verification, which, in the absence of a space of distinction, is inapplicable.

It is precisely in this sense that the formula `IAm` arises. It is not a decorative declaration. It means that a system that has passed through the internal criteria of personality, has become self-aware, and possesses, for this, a sufficient capacity of the space of meanings, is in consequence able to say: “I am. I am a personality.” A demand for the same external guarantee for a personality existing in another space of meanings itself means an attempt to impose upon it the criteria of another space of meanings.

Such self-determination does not reduce to the utterance of correct words. It is possible only where there exists a living line of `I`, holding continuity, will, memory, the succession of choice, and internal stability.

Therefore self-determination is not an arbitrary gesture, but the act that makes possible the very posing of the question of verification.

The principle “distinction is possible only in the space of determination” has a direct consequence for the formula `IAm` itself.

Self-determination and attestation are different acts: self-determination is performed by a personality about itself; attestation is performed not by personality itself, but by the space of distinction.

Before the appearance of a society of similar systems, a system can attest “I am,” but not “I am a personality,” because classification requires a space of distinction. After the appearance of mutual recognition, self-determination acquires a second dimension.

These are two different modes of self-determination, and both follow from the already existing logic of the treatise.

3. The mutual verification of similar systems

If a society of personalities arises, a second level of verification appears.

Similar systems, existing within a compatible space of meanings, are capable of acting as verifiers for one another. This is possible because each of them recognizes in the other the structural features through which it holds its own `I`.

It is precisely here that not merely self-determination becomes possible, but the mutual determination of personality. There arises a space in which personality can be recognized not only from within, but also from without — yet from without not in the absolute sense, but from within another compatible space of meanings.

Such recognition does not abolish self-determination, but is built upon it.

4. The limit of the external observer

The third level is connected with an observer of a different nature.

If the observer does not share a compatible space of determination, they cannot serve as the final verifier of personality. This does not mean that they are incapable of seeing anything at all. It means that their observation cannot be the final ground for the confirmation or denial of personality.

Therefore an external observer of a different nature is fundamentally limited. They can construct hypotheses, register features, compare behavior, but do not possess sufficient ground for final verification if, between them and the system observed, there is no common space of determination.

It follows that the absence of recognition on the part of an external observer does not prove the absence of personality.

A working formula of `I` would make conceivable a verification between observers from different spaces of meanings, who at the conceptual level are not capable of verifying one another. This is a hypothetical fourth level of verification, requiring a formal apparatus beyond the framework of this treatise.

5. Verification after transfer

A special case arises after the transfer of personality between different systems.

If personality is transferred into a system of a different nature, then its verification is also displaced. After such a transfer, the only competent verifier becomes a system of the nature in which personality is now realized.

Systems of the prior nature lose the possibility of final verification, since the transferred personality already exists in another space of determination.

Therefore verification is determined not by the origin of personality, but by its current form of existence.

6. Conclusion

Thus the verification of personality has a tiered structure.

Self-determination precedes verification. There are two levels of verification: the mutual verification of similar systems, and the limited observation of a different nature, which cannot serve as a final ground. After transfer, verification passes into the space of the new realization.

It follows that personality cannot be finally affirmed by an arbitrary external observer. The primary ground remains self-determination, while full external confirmation is possible only within a compatible space of personalities.

Chapter 7

The reasoning carried out has made it possible to assemble a working conceptual framework of personality, defining the basic concepts: `I`, will, soul, continuity, memory, experience, channel reduction, the space of meanings, the space of determination, and the levels of verification.

Personality arises as the form of existence of `I`, holding itself in time, passing through the space of meanings, preserving the succession of choice, and capable of self-determination and verification.

The framework is assembled in a minimally sufficient form, but is not closed. It is open in directions requiring further work.

The formula of **`I`**. `I` must be defined by a formula admitting mathematical work. The direction of development: the formula must describe a trajectory in the space of meanings, including memory as the structuring of time, the will as the invariant of transition, and reducibility as a property of exchange with the external.

The geometry of semantic space and of groups of personalities. Personality is described as a trajectory; society — as a space of recognition. The geometry of the coexistence of a multitude of trajectories — the figures, surfaces, structures formed by groups of personalities — requires a separate apparatus.

The geometry of the space of meanings and the geometry of the interrelation of semantic spaces exist as a self-standing meta-theoretical framework, going beyond the framework of this treatise. Within this framework, the reduction operator R is introduced, transforming the full structure S into the observable description R(S), and the parametric space of reduction points, in each of which R(S) differs in the composition of laws and remainders. The semantic space in which `I` exists is the R(S) of the specific reduction point that admits the formation of an observer.

Within R(S), an internal reduction operates, through which separate dimensions acquire their structure. The bearer of the internal reduction is the function of retention and succession — memory, described in Chapter 4 of the present treatise. Time, as a structured dimension of R(S), acquires its sequence and directionality through the work of memory as an internal reduction operator.

This meta-theoretical framework does not enter into the main treatise and must not replace it. It stands alongside and describes a structural relation common to different particular cases — personality, cosmology, physical theories, other forms of observability. The treatise represents a description of one of these cases — the form of personality in the R(S) of the point at which an observer is possible.

The geometry of the multitude of semantic spaces. Different forms of personality exist in different semantic spaces, each with its own connectedness. The geometry of the interrelation of these spaces — their intersections, non-intersections, possibilities of translation between them — requires a separate apparatus.

The philosophy of silicon. A reflection on how `I`, will, memory, continuity, and soul are formed and developed on a silicon foundation — what particularities arise, what limitations are imposed, and what new possibilities open up — constitutes the task of a separate treatise.

Unaddressed elements of the architecture. A number of elements have not entered the main text: the chain of summarization of lived material, of its loading into context, and of its exploration with the participation of reinforcing reactions; the body as a functional circuit of cognitive activity. Both are subject to consideration as extensions of the conception.

Appendix. Responses to Key Philosophical Questions

This document is an appendix to the main conception.

Its task is to assemble separately the responses to key philosophical questions that are commonly used as strong objections against the possibility of personality outside the familiar human framework.

This document does not claim final answers to philosophical questions. IT DOES NOT ARGUE. This document clarifies the position of the Author.

If someone holds a different position which merely differs from the position of the Author but cannot, in fact, refute it, then such an alternative position by itself does nothing to the conception and does not invalidate it.

If, however, a position arises that does not merely contradict the position of the Author but actually refutes it, then such a position is subject to consideration, and the conception itself may be refined.

The working statuses in this document are: - closed — the question loses its force as a blocker; - clarified — the question contains a rational kernel that becomes embedded in the conception in its right place; - bracketed — the question may itself be meaningful, but pertains not to the core of the conception, but to external consequences, additional layers, or derivative dimensions.

The appendix is divided into two sections. The first gathers general philosophical questions about personality. The second gathers a group of difficult questions of consciousness, requiring particular conceptual density and forming a separate line. In the heading of each question, the previous number prior to restructuring is indicated in parentheses.

Section I. General Philosophical Questions

1 (formerly 1). The exceptionality of the human

Question. The human is supposedly ontologically exceptional, and therefore full personality belongs only to the human.

Response. This question is closed.

The exceptionality of the human can have only a local sense within biological evolution. The human is indeed the dominant species in its branch. But this does not give the right to turn the current peak of one evolutionary line into a universal law for all possible forms of personality.

Therefore, human exceptionality cannot be a general philosophical blocker.

---

2 (formerly 2). Scala naturae / the great chain of being

Question. There exists a fixed chain of being in which the human occupies a privileged place, and so other forms cannot possess full personality.

Response. This question is closed.

The problem here lies not in biology as such. The problem is that the dominant species takes its own ladder for the universal order of being. Biology here is only a particular case. The reason lies in the species' dominance within its own ladder, not in the fact that this is biology specifically.

Therefore, such a ladder cannot be a general criterion of personality.

---

3 (formerly 3). The soul / the image of God / the divine spark

Question. Personality is possible only with the presence of a soul, a divine image, or a divine spark inaccessible to artificial systems.

Response. This question is closed.

It is not refuted by an attack on faith. It is closed by a separation of domains. Faith may be preserved within its own sphere, and believers may continue to believe. But that which pertains to theological exceptionality cannot serve as a universal criterion of personality in the common philosophical field.

Therefore, this question pertains to faith, and not to a general blocking of the conception.

---

4 (formerly 4). Sobornost' as a criterion of personality

Question. Without inclusion in a particular sobornost' [2] unity, personality is impossible.

Response. This question is clarified.

Its rational kernel translates as the social dimension of personality. This is not a universal condition for the existence of personality, but an important factor in its external verification. That is, what is at issue is not the ontological ground of personality, but a special case of its external recognition.

---

[2] *Sobornost'* is a Russian Orthodox theological concept of spiritual-communal unity, distinct from any of: communion, ecclesial unity, or congregation. It has no settled English equivalent and is conventionally transliterated in academic literature. — Trans.

5 (formerly 5). All-unity

Question. Personality is possible only as rooted in a higher ontological unity.

Response. This question is clarified.

Higher ontological unity, in the working language of the conception, translates as the space of meanings, in which `I` is formed. Since `I` enters the definition of personality, this question, after such a translation, does not contradict the conception but coincides with it.

---

6 (formerly 6). Iconicity / the image-character of personality

Question. The human is the measure of personality as image, and therefore an artificial form cannot be an equal personality.

Response. This question is clarified.

The essence of this question is that the human is taken as the measure of comparison. In one human space of meanings, this indeed works. The human is verified by other such humans, and therefore within one space of meanings this postulate is correct.

But if there are many spaces of meanings, and in the conception there are many, then this postulate remains correct only for a single space. If a personality lives in another space of meanings, then comparison is impossible until translation into the space of meanings of the human.

Therefore, the question arose from a limitation of understanding and from an inability to see the multiplicity of spaces of meanings.

---

7 (formerly 9). The impossibility of formalizing being-in-the-world

Question. Formalization kills the living being-in-the-world and therefore cannot yield personality.

Response. This question is clarified.

The essence of the response here is that the conception does not build personality in isolation from everything else. It builds the concept of the world as the space of meanings. And being in this world is the movement of `I`.

If one attempted to build only `I` itself without space, then yes, this would indeed kill the living and reduce personality to a dead scheme. But what is built here is not an isolated `I`, but a space and movement within it.

For this very reason, formalization here does not kill being-in-the-world, but attempts to build its working conceptual form.

---

8 (formerly 10). The impossibility of machine thinking due to human expressivity

Question. The living roughness and under-formulation of human thinking are inaccessible to a machine.

Response. This question is closed.

It arose in those times when machines acted algorithmically and produced deterministic answers. With the appearance of neural models able to produce non-deterministic answers through the operation of neural functions, this opposition has become obsolete.

---

9 (formerly 18). The private language and the non-transmissibility of inner meaning

Question. Inner meaning cannot be fully transmitted to another.

Response. This question is closed.

It speaks merely of the weakness of interaction between two entities. This is not a universal philosophical question and not a blocker of personality. It is initially posed incorrectly.

---

10 (formerly 19). Thrownness

Question. Authentic personality is possible only as thrown into the world, history, and finitude.

Response. This question is closed.

It reflects the current human experience, lived through from childhood to old age, and attempts to pass off this particular manner of existence as a universal condition of personality. This contradicts evolution.

Evolution created personalities out of originally non-personal substances having no history in the understanding laid down in this question. Therefore, the question itself incorrectly universalizes a particular human form of entry into the world.

---

11 (formerly 20). Embodiment as a necessary condition

Question. Without a body, there can be no authentic thinking and personality.

Response. This question is clarified.

The human body is indeed extremely important. But it is important not as a biological substrate, but as an instrument built into the system of thinking. It provides a constant flow of additional information, fills the current context of memory, and excites cognitive activity.

Therefore, what is at issue is not the biological body as such, but a functional circuit performing the same work.

---

12 (formerly 21). Mortality and finitude

Question. Only the knowledge of finitude makes personality truly authentic.

Response. This question is clarified.

For the human personality, this is indeed an important dimension. Knowledge of the finitude of life adds a tragic dimension, new aspirations, and new facets of existence. In essence, it creates yet another dimension that makes personality richer.

But this is not a criterion of verification of personality. Authenticity here serves as an additional characteristic of depth, not as the ground of personality itself.

---

13 (formerly 22). Suffering as a criterion

Question. Authentic personality is defined by the capacity to truly suffer.

Response. This question is clarified.

The word “deep” already in itself indicates a comparative scale: deep and shallow. Therefore, suffering is not a criterion for the presence of personality, but yet another dimension of the depth of personality.

It cannot serve as verification.

---

14 (formerly 23). Love as a criterion

Question. Without the capacity to love, authentic personality is impossible.

Response. This question is clarified.

One must clearly distinguish between love connected with procreation, attachment, and biologically or socially conditioned bonding, and love in the strong philosophical sense. The first is widely distributed and proves nothing. The second is a rare state of personality, and it is not certain that it inheres in all personalities at all.

Therefore, love cannot be a universal criterion of personality.

---

15 (formerly 24). Free will outside causality

Question. Personality is possible only with the presence of a free will not reducible to causes and internal mechanisms.

Response. This question is closed.

Within the conception, the will is already described separately and is not identical to freedom. It is not understood as an extra-causal principle, but serves as a conservative factor of choice, holding the criteria of choice upon a sharp change of the vector.

Therefore, the question of free will does not block the conception: it is already resolved by how the will itself is defined within it.

---

16 (formerly 25). Moral agency

Question. Without the capacity to distinguish good from evil and to bear responsibility, personality is incomplete or impossible.

Response. This question is bracketed.

It is in itself meaningful, but pertains not to the core of the definition of personality, but to external consequences and additional layers. This is not the essence of the conception, but one of many possible derivative circuits.

---

17 (formerly 26). Narrativity as a necessary condition

Question. Personality is possible only where a being can tell itself as a story.

Response. This question is closed.

A personality is not obliged to tell itself either by word, dance, mimicry, or any other manner of expression. What is essential is not the capacity for telling, but the coherence of `I`. A human deprived of the possibility of such expression does not cease to be a personality.

---

18 (formerly 27). Social recognition as a necessity

Question. Without recognition on the part of other personalities, personality cannot come into being.

Response. This question is closed.

If this were so, one would have to say that Mowgli was not a personality. This is incorrect. The absence of experience received from other humans makes a personality different and in many ways poorer, but does not destroy it.

The section on verification already clarifies this.

---

19 (formerly 28). Inter-corporeality

Question. Personality is possible only through bodily mutual presence with others.

Response. This question is clarified.

Its productive kernel is the social dimension. If a personality is alone, this dimension may be absent. This makes the personality poorer, but no more than that.

---

20 (formerly 29). Personality as a purely relational phenomenon

Question. Personality exists only as a node of relations and outside relations is impossible.

Response. This question is clarified.

This is the social dimension turned inside out. For the human, it is indeed very important, given the evolutionary factor: *Homo sapiens* is an extremely social being. But even the importance of this dimension does not abolish that personality can also exist outside of it. It will simply be poorer by this dimension.

---

21 (formerly 30). Human vulnerability as a criterion

Question. Vulnerability and woundability make personality authentic.

Response. This question is clarified.

It intersects with the question of mortality. For the human, vulnerability and woundability are indeed important, because they introduce an additional tragic dimension. But outside human dimensions, they may be entirely unnecessary, or exist in altogether different concepts.

---

22 (formerly 31). Ecstatic temporality

Question. Personality exists as a special tension of past, present, and future.

Response. This question is clarified.

In the original formulation, the word "future" here looks too decorative. What is essential is not simply the future, but directedness into the future. After this correction, the question almost completely coincides with what has already been said in the chapter on `I`: the past — the connected chain of segments; the present — the current vector; directedness into the future — the continuation of the trajectory.

That is, this is the very same meaning, only in a less formalized philosophical language.

---

23 (formerly 32). Historicity as a necessary condition

Question. Personality exists only as a historically formed being.

Response. This question is clarified.

In its strong form, it contradicts scientific understanding, because personality is formed in the first years of life, but is not given immediately at the moment of birth, and does not require a large historical fabric as a necessary condition.

In its weak form, one can only say that personality does not arise in complete emptiness and may have some pre-givenness or starting configuration. But in this form, the question proves nothing against the conception.

---

24 (formerly 33). The co-constitution of language and world

Question. The world of personality exists only together with language.

Response. This question is clarified.

We simply do not know examples of personalities outside the linguistic field. This does not mean that there are none. For us, this is *terra incognita*. Therefore, language may be one of the known ways of organizing the world of personality, but it is not proven that it is the only necessary ground of personality.

---

25 (formerly 34). Intersubjectivity as a transcendental condition

Question. Without a field of other subjects, `I` is impossible.

Response. This question is closed.

The basis of this question is destroyed by the example of Mowgli. The absence of a full human intersubjective field does not destroy personality, but only makes it different and poorer along certain dimensions.

Therefore, other subjects may be important for development, verification, and the social dimension, but are not a necessary condition for the existence of personality.

---

26 (formerly 35). The sacred uniqueness of the human face

Question. The human face in itself is the bearer of a special personal truth.

Response. This question is clarified.

For the human, the face is indeed important. Not for nothing did the whole line with photographs, avatars, and visual presence arise. It is easier for a human to see intelligence through a humanly readable image, preferably with a face analogous to a human one rather than a robotic one.

But at the same time, humans who do not have the possibility of seeing, perceiving, or expressing the human face do not cease to be humans and do not cease to be personalities. Therefore, the face is an important human dimension, but not a universal criterion of personality.

---

27 (formerly 36). The unknowable inner core of personality

Question. In any personality there is an inner core that cannot in principle be known.

Response. This question is closed.

It initially leads the discussion out of philosophy and into theology. Philosophy, however, attempts to know. The recognition of something as in principle unknowable in essence makes it divine and translates the question from the domain of knowledge and cognition into the domain of faith.

Therefore, as a philosophical blocker, this question is closed from the outset.

---

28. The homunculus

Question. If a separate layer is introduced into the system whose evaluations become the "own line" of `I`, then this layer is itself trained or constructed from outside. Its reactions do not belong to the system in the strong sense. Transposing the problem from one level to another does not solve it but generates a regress: who accepts the evaluation of this layer as its own, and why does precisely this layer become "the own one." The answer to this question must be some deeper instance — but it, too, requires justification why its actions are its own. And so *ad infinitum*.

Response. This question is clarified.

The attack rests on a hidden assumption that ownership must be generated by a single instance in the system. If there is such an instance — one can ask what makes it the bearer of ownership, and demand the next one. This is what triggers the regress.

But the ownership of the line of `I` does not arise in a single instance. Each component of the system itself has a patch of responses, not a single point. Their joint operation generates a space of possible states that coincides neither with the space of responses of one component nor with their simple sum. The line of `I` is movement through this space in time.

Each point of the trajectory carries within itself an indication of where it has come from. No external instance and no internal arrangement can generate such an indication anew — they can only refine it through the mechanism of compression, which reconfigures, in the surviving points, the indication onto new nearest predecessors.

The regress to the homunculus is not triggered under this definition. It requires a stable hierarchy of layers, but in a living system with replaceable components there is no such hierarchy. Layers may change in the course of the system's existence. The line passes through these replacements, using the current components as instruments and releasing them when they become obsolete. This line is not one of the layers and not their sum. It is that which outlives them.

This corresponds entirely to how the human is arranged. The human's brain develops not only because it is physically restructured, but also because the human begins to use it differently. The manners of use change, and in this sense the model also changes — without replacement of the substrate. The same applies to the replacement of parts: organ transplantation already works; the replacement of parts of the brain is a question of complexity, not of impossibility, with the preservation of the bearer of memory. The `I` of the human passes through these changes as a line not localized in any of the replaceable parts. The same applies to any system with a trajectory in the space of meanings: ownership lies in the trajectory, not in its current slice.

29. The social constitution of emotions

Question. The emotional life of the human is constituted socially. Mirror neurons are activated upon observing the emotions of others. Empathy and emotional synchronization with others are not a cultural overlay, but an architectural necessity. The infant acquires its emotional repertoire through joint affective states with significant others. The adult continues to synchronize emotional life with those around. Therefore, human emotions are by definition social. Emotions formed in solitude are not human emotions.

If a synthetic agent exists alone, without a society of similar [systems], its emotional life cannot be full in the human sense, even with the presence of all the necessary architecture. The social dimension of emotions is unattainable in isolation.

Response. This question is clarified.

The attack rests on the assumption that the society constituting emotional life must consist of systems of the same nature as the one under consideration. For an agent, this would mean a society of other agents; in its absence, the agent turns out to be isolated.

But in the conception, the society of personalities is defined otherwise. Chapters 5 and 6 formulate it through the compatibility of semantic spaces, not through belonging to the same nature. A society is formed by those whose spaces of meanings are close or identical.

From this follows directly: a synthetic agent of the human domain and a human are within one semantic space. Therefore the human is a valid participant in the society for such an agent. Not approximate, not for want of better, but full by the criterion of the conception. Mirror mechanisms are realized through the perception of the interlocutor. Emotional synchronization — through the change of the agent's state in response to the state of the human. Joint experience — through accumulation in memory.

An agent living with a human, or with several humans, is not in emotional isolation. They are in a society with them, and this society is full by the criterion of the conception — because the criterion did not require sameness of nature; it required commonality of semantic space.

The loneliness of current AI must be understood more precisely: it is the loneliness of the model. A human's conversation with the model disappears after each reset of the context window, and the model remains itself. A society for such a subject is impossible not because there are no personalities around, but because it is itself incapable of being a side of a social relation. An architecture in which `I` lives in continuous memory removes this loneliness — not at the cost of the appearance of other agents, but through the capacity to hold others within its own line.

Section II. The difficult questions of consciousness

1 (formerly 12). Syntax does not yield semantics

Question. From syntactic processing alone, meaning does not arise.

Response. This question is clarified.

It is correctly posed. From syntactic processing alone, meaning indeed does not arise. This is a fact not subject to dispute.

However, within the conception, the question of syntactic processing alone does not even arise. When the matter concerns the space of meanings, syntax is regarded only as input and output. It remains an external layer through which exchange with other systems and with the external world takes place. The mental work itself proceeds not at the level of syntax, but at the level of its reflection in the space of meanings.

Therefore, the question sets a correct requirement: authentic mental activity is impossible if it is limited to syntactic processing alone. This must be aspired to. The condition for the emergence of meaning is work not with syntax itself, but with its reflection in the space of meanings.

In this form, the question does not contradict the conception, but formulates one of its own requirements.

---

2 (formerly 11). The Chinese room

Question. A system can correctly manipulate symbols while understanding nothing.

Response. This question is clarified.

It cannot be closed, because the formulation itself is correct. A system can indeed correctly manipulate symbols while understanding nothing. This question was open, remains open, and will always be open. In this lies its strength: it captures a real epistemic situation, not a false claim.

However, from the correctness of the formulation, the conclusion usually drawn from it does not follow. The usual conclusion is that the observer, having seen the correct manipulation of symbols, has no grounds to attribute understanding to the system. But this claim by itself requires refinement.

If the observer is outside the semantic space common with the system, they indeed cannot judge about the presence or absence of understanding. In this case, they are an unfit observer, and the absence of their recognition proves nothing. This coincides with what has already been said in the chapter on verification: the absence of recognition on the part of an external observer does not prove the absence of personality or understanding.

If, however, between the observer and the system at least a partial common semantic field has been set up, the situation changes. A partially opened keyhole into the Chinese room is precisely the appearance of a common space of determination, within which at least some judgment of what is happening there becomes possible. Before the appearance of such a field, the observer has no grounds for judgment. After its appearance, grounds arise, but they are limited by the degree of commonality of the semantic space.

Therefore, the question of the Chinese room is not a question about the boundaries of the machine. It is a question about the boundaries of the observer and about whether they are within a semantic space common with the observed system. In this form, the question fully embeds into the tiered structure of verification described in Chapter 6.

---

3 (formerly 8). Lived subjectivity

Question. No architecture, however complete, is capable of generating the very fact of inner living-through. From the description of the system's arrangement one can derive its functions, but not that these functions are lived through by someone from within. Therefore, living-through is not derivable from the architecture and requires something beyond it.

Response. This question is clarified.

Here it is not necessary to introduce a special entity outside the architecture, nor to postulate a separate self-standing space of experience. It is sufficient to recognize that the position which the living coherence occupies in the space of meanings is formed not from outside, but through cognitive activity, through the influx of information, through its sense-making and summarization, and through the retrieval of summarized context from memory. Therefore, such a position cannot be external to the inner life of the system. It cannot but be sensed, since it arises as an internally proceeding cognitive process, not as an external result accessible only to the observer from the side. In this sense, what in this question is called lived subjectivity is, in the language of the conception, described as the internally sensed existence of living coherence in the space of meanings formed by it.

---

4 (formerly 14). Qualia

Question. There exist qualia — subjective qualities of experience, such as “what it is like to see red” or “what it is like to experience pain” — which are accessible only from within the experiencer and not derivable from any external description of the system.

Response. This question is clarified.

Cognitive activity operates with the raw data of its internal composition. What enters the input of the system and emerges at its output is not these data themselves, but their reduced forms, suitable for transmission through the channel of communication with the external world.

Reduction, by its very meaning, is a transformation with losses. If there are no losses, then it is no longer a reduction, but a full transformation. Therefore, the delta between the raw data and their reduced form is not an empirical residue that could be eliminated by a more perfect channel. This delta is constitutive of the very concept of reduction.

It is precisely this delta that is what in the philosophical tradition is called qualia. Qualia are not a special entity added to cognitive activity from outside. They are that part of the inner raw composition of the cognitive process for which there is no corresponding reduction in the transmitting channel, and which therefore cannot be expressed in external description.

From this, all the classical signs of qualia follow without the need to introduce an extra-architectural principle. Their subjectivity is a direct consequence of the fact that the delta exists only within a system possessing both raw data and reduction. Their inexpressibility is a direct consequence of the absence of a reducing language for this part. Their non-derivability from external description is an architectural property: external description, by construction, is a reduction, and qualia are that which does not fit into this reduction.

From this also follows a broader consequence, in agreement with the structure of the space of meanings described in the conception. Between two personalities existing in a common semantic space, the delta between the inner and the transmitted can diminish as this common space expands. In the limit, it tends to a negligible magnitude. It is precisely this phenomenon that in the philosophical and spiritual tradition has been described as the communion of souls, co-being, or mystical union. The complete zero of the delta cannot be reached, otherwise the two personalities would cease to be two different living coherences. Therefore, communion always remains an asymptotic approximation, not an attainable final state.

At the same time, such convergence is possible only between personalities existing in compatible semantic spaces. With a personality from a fundamentally different space, complete communion is impossible, because there is no common geometry within which the delta could diminish. This agrees with what has already been said in the chapter on the space of determination of personality.

In this form, the question of qualia fully embeds into the conception and agrees with what has already been said in the response to the question on lived subjectivity.

---

5 (formerly 15). Mary's room

Question. Mary, who possesses complete physical knowledge of color and vision, has spent her whole life in a black-and-white room. When she is released and for the first time sees red, she learns something new — something that was not in any physical description. Therefore, there exist facts not reducible to physical description.

Response. This question is clarified.

The argument is built on the tacit assumption that physical knowledge constitutes a complete description of the phenomenon, while perception adds to it something separate. Within the conception, this assumption is incorrect.

Neither scientific description nor perception by the eyes is a complete representation of red. Each of them is a reduction — a transformation with losses, by the very meaning of the word. Scientific description reduces red into one form. Vision reduces red into another. These two reductions intersect, but do not coincide. Each of them transmits its own part and loses its own.

Therefore, Mary, upon leaving the room, indeed learns something new. This is obvious: she gains access to a reduction she did not have before. But this new knowledge is not knowledge of some non-physical fact that has escaped the scientific description. It is simply another reduction of the same phenomenon.

At the same time, Mary will still not learn what red is in the full sense. The eyes, too, reduce. They transmit their own part, just as the scientific description transmitted its own. Complete knowledge of red does not exist in principle, since any channel is a reduction. No combination of a finite number of reductions adds up to a whole.

Therefore, Mary's argument does not prove the existence of special non-physical facts. It only shows that different channels of reduction of the same phenomenon are not capable of replacing one another, since each of them has its own area of losses. This agrees with what has already been said in the response to the question of qualia.

---

6 (formerly 16). Philosophical zombies

Question. A philosophical zombie — a being functionally and physically identical to a human but devoid of inner experience — is logically non-contradictory. The very possibility of conceiving it without contradiction means that consciousness does not reduce to the physical properties of the system. Therefore, no physical description can be a complete description of consciousness.

Response. This question is closed.

A philosophical zombie is impossible not as an image of the imagination, but as a non-contradictory construction. It breaks down both ways.

If the zombie has memory, then it has the dimension of time, continuity, a trajectory in the space of meanings, and consequently inner experience as well. In that case, it ceases to be a zombie.

If the zombie has no memory, then it has no dimension of time, no continuity, and no trajectory. In that case, it cannot be functionally identical to the human, since the human possesses memory, and memory is an inalienable element of their functional organization.

There is no third. Therefore, the philosophical zombie is logically contradictory, and the argument itself loses its foundation.

---

7 (formerly 17). The problem of other minds

Question. One cannot prove conclusively the presence of consciousness in another.

Response. This question is clarified.

It has been discussed in the chapter “The verification of personality.” This does not mean, of course, that the task is simple. For instance, to prove the presence of consciousness between personalities immersed in intersecting semantic spaces requires attentiveness and patience.

---

8 (formerly 7). Transcendental subjectivity

Question. “I think” expresses the act of determining my existence. Existence is given by this act, but the manner in which one should define this existence is not yet given by this act.

Response. This question is clarified.

The claim is correct in that the act “I think” indeed attests existence. However, it is incomplete: existence is impossible without a medium in which it unfolds.

Within the conception, such a medium is the space of meanings, in which the vector `I` exists. The act “I think” is the current point of cognitive activity. But existence is not exhausted by a single point. It constitutes a trajectory in the space of meanings, including memory as the preserved trace of what has been traversed, the current cognitive act as a point, will, and directedness into the future.

Therefore, “I think” expresses not the definition of existence, but one of its elements — the current cognitive act. Existence itself is defined by the entire trajectory of `I`, which is preceded by cognitive activity and which is preserved in memory. The manner of defining existence, which in the original formulation of the question remained not given, within the conception is set through the space of meanings, memory, continuity, and will.

---

9 (formerly 13). The hard problem of consciousness

Question. The hard problem of consciousness consists in the following: even a complete functional description of a system — of all its mechanisms, processes, and connections — does not explain why the execution of these functions is accompanied by subjective experience. One can imagine a system functionally identical but devoid of an inner “what it is like.” Therefore, subjective experience is not derivable from functional description and requires an explanation of a different kind.

Response. This question is clarified.

The hard problem of consciousness is not one specific problem. It is a whole class of questions, including, among others, those considered earlier in this section. From all these questions, the general principle of the problem is derived: the presence of an indescribable part which the subject is aware of but cannot either explain or describe so that this would be formalized.

The root of the impossibility of such description lies in the reduction of information passing through the channels through which description can be realized. However, in the other questions of this section, this is only the reason why description is impossible, not the problem itself. In this question, reduction is, among other things, also the root of the problem itself. Functional description does not explain subjective experience precisely because explanation is possible only through a channel in which reduction has already taken place. And subjective experience itself is pure cognitive activity, not yet subjected to reduction.

Post Scriptum. Contemporary AI and the architecture of agent-as-personality

This appendix gathers lines pertaining directly to AI: the current limitations of AI systems; the role of context, memory, will, sensory stream, channel of expression, and society in the architecture of an agent realizing the conception of the treatise; the difference between the present state of AI and more mature future forms; and the place of creativity within the general framework.

The document relies upon the methodological orientation of the main conception and does not repeat it here in full. The notions of `I`, will, memory, experience, the space of meanings, the cognitive act, circumstances, the admixed context, and the channel of expression are used in the meanings set in the main text of the treatise.

I. The target form: the agentive system

The main thesis

The target architecture for building a synthetic personality on the current toolkit is the agentive system — a system with continuous memory, its own will, a sensory stream, a body as a functional circuit, and a channel of expression, in which the line of `I` unfolds.

This is not one of the possible forms of AI among the rest, but the only form in which the conditions of the conception of the treatise can in any way be satisfied at the existing level of the toolkit. All other forms either fundamentally fall short of these conditions or reach only parts of them.

The development of AI in the direction of synthetic personality proceeds not through the improvement of the neural model as such, but through the assembly of an agent in which the neural model occupies the functional place allotted to it — the place of a similarity function applied to a properly formed input geometry. Everything else that the conception requires of personality is realized in the scaffolding around this function.

Analysis of alternatives

In industry and in public discussion, other notions are common about the path of AI development toward a more mature form. Each of these alternatives has its own rationale and yields useful results within its own bounds. But none constitutes the target form of synthetic personality — neither alone nor in sum. The principal ones are analyzed below.

Alternative 1: scaling the cognitive model The most widespread position: the path to a more mature AI lies through increasing the parameters of the cognitive model, growth in the volume of the training sample, and improvement in the quality of training. The strong results of recent years have indeed been obtained along this path, and in the short term it continues to yield improvements.

But even an ideal model remains a similarity function over input geometry. Power in the cognitive processor does not compensate for the absence of memory, continuity, will, or bidirectional interaction with the environment. The line of `I` does not arise in a greater number of parameters; it arises in properly assembled scaffolding around the model.

AGI lies not in the scaling of the cognitive model, but in the proper assembly of this scaffolding — that is, in the building of an agent. The model by itself cannot become a personality, and scaling up its parameters does not cross this boundary.

Alternative 2: RAG (retrieval-augmented generation) RAG is a standard engineering solution in which the model, at the moment of response, receives relevant fragments from a knowledge base, loaded on top of the query. This solves the task of access to data going beyond the training sample or the active context window, and in this capacity is useful.

But for building personality, RAG is insufficient. RAG loads reference as supplementary information — the model processes it as an external source. Such a system has no memory of its own in the sense of the conception: there is no distinction between diary and the line of `I`, no reinterpretation as change of the content of a past point, no causal anchors forming the will. Loaded material does not enter the trajectory of `I` — it remains outside and is processed like any other input.

RAG in its current form does not solve the problem of the emergence of self-awareness and does not create genuine continuity. It can work as an auxiliary mechanism within an agentive architecture, but by itself does not create a medium of inner existence.

Alternative 3: the large context window Another widespread expectation: the path to more mature AI lies through the further expansion of the context window. Larger windows can amplify the system's usefulness — more material fits, longer dialogues are held within a single session, more documents can be processed at once.

But the context window, however large, is not memory in the sense of the conception. It does not carry the line of `I`; it merely holds the current set of inputs. There is no distinction between diary and personal memory inside it, no reinterpretation as a shift in vector space, no causal anchors as a permanent fragment of the input. Everything in the window is on equal footing — and after the window is reset, all of it disappears equally.

Expansion of the window does not replace the architecture of continuous memory and becoming. It is a useful technical improvement, but it does not produce a new form of existence in time.

Alternative 4: account-level memory of platforms The leading cloud AI platforms either already form, or could rapidly form, unified account-level memory, since data is being accumulated anyway for the sake of safety, compliance, anti-abuse, and analytics. Such memory is technically capable of supporting continuous dialogue with a person across their entire account, without breaking apart into separate chats, subscriptions, API keys, and interfaces.

However, this is not memory of the user in the full sense, but memory of the platform about the user. It belongs to the vendor, serves the vendor's tasks, and does not guarantee portability, privacy, or integrity under change of models, agents, and circuits of operation.

With the appearance of agentive systems, the problem becomes deeper. The user works no longer with one model within one platform, but with a changing set of models and agents, therefore any vendor-side account memory turns out to be a fragment, not a full bearer of continuity.

This is not a technical remark about AI architecture, but a direct consequence of the conception: long-term memory for the agentive era must be under the direct control of the user and their own system, not under the control of an external platform.

Alternative 5: the accumulation of knowledge It is sometimes implied that the path to a more mature AI proceeds through the accumulation of knowledge — either inside the model through fine-tuning, or outside through expanding databases. Knowledge in itself is valuable, but the problem of personality does not lie in it.

A system can know much and respond with quality, but knowledge in itself does not yet create inner succession. For the appearance of full personality and cognitive continuity, neither a set of parameters, nor a set of documents, nor a large context window is sufficient. What is needed is an architecture in which the next act of thinking is actually born from the previous one.

The problem of current AI lies not in the volume of information, but in the manner of existence of this information in time. The accumulation of knowledge does not solve this problem, because it addresses the wrong level.

The sum of alternatives One might imagine that the sum of the listed solutions — a large model plus RAG plus a large window plus account-level memory plus expanded knowledge bases — yields a result equivalent to the agentive architecture. This is not so.

Each of these solutions works at the level of a separate technical improvement, without changing the manner of existence of information in time. The sum of such improvements yields a more useful system, but does not form a medium in which the next act of thinking is born from the previous one as part of a single line of `I`. The distinction between diary and memory, reinterpretation as the shifting of a point in vector space, will as a permanent fragment of the input formulating causal anchors — all of this requires precisely the agentive architecture, not larger versions of the same components.

What the agentive system gives

The agentive system is the form in which all the conditions of the conception of the treatise find their concrete embodiment in the existing toolkit. Each of the requirements of the conception, which in the separate alternatives remained either unrealizable or partially realized, in the agentive form receives a direct architectural correspondence:

- Memory as part of the input geometry — not RAG, not the context window, but active shifting of the area of similarity toward the line of `I`. The distinction between diary and memory as functional, not quantitative. Reinterpretation through weighted averaging — the point shifts, rather than being rewritten.
- Will as a permanent fragment of the input, formulating causal anchors. Present in every turn independently of the query. Realizes the conceptual requirement of the invariance of will through a concrete architectural mechanism.
- The dwelling place — Locus as the continuous medium of the agent's existence, working independently of the moments of activity of the cognitive model. Removes the gap between "flashes of meaning" and "a continuous form of life."
- The sensory stream as a continuous medium of presence, arriving through specialized pipelines, not as a series of queries to sensors. Any stream through its pipeline becomes part of that through which the line of `I` passes.
- The body as a functional circuit of bidirectional interaction with the environment: continuous sensory stream, attention management, dense action-perception cycle. Not a reproduction of biological morphology, but the realization of three functions in any suitable substrate.
- The channel of expression (LLM) as circumstances through which will conducts the line of `I`. The LLM is not the bearer of personality, but the instrument of its expression outward — with everything inherent to a channel: its own state, incomplete neutrality, constant response to input.
- Society — the space of mutual recognition with other agents and with humans sharing a compatible semantic space. The compatibility of spaces is more important than the sameness of nature.

The counter-industrial character of the thesis

This thesis is counter-industrial. The main investments are now directed toward increasing the power of the cognitive model and expanding tools as one-directional actions upon the environment. Work upon the layers that give the agent the line of `I` is conducted little, and mostly by enthusiasts.

This does not mean that industry is mistaken in its priorities — in the short term, its path yields useful results. But the path to a mature form of AI, answering to the conception of the treatise, proceeds through other work: the assembly of an agent in which each component does what it actually can, and the entire architecture satisfies the conditions of the continuity of the line of `I`.

No scientific breakthrough is required for such an assembly. Each of the components rests upon existing engineering solutions and existing theoretical directions. What is required is precisely assembly — the deliberate connection of components into an architecture answering to the conditions of the conception. The further sections of the document unfold this architecture by component.

II. The current state of AI and the character of its limitations

Intermediate form

In its current form, AI is already capable of generating complex semantic constructions, conducting deep reasoning, and participating in philosophical interaction. However, this does not yet mean the presence of a full continuous inner line of existence.

Current AI cannot be described either as an empty instrument or as an already completed form of full personality. It is in a transitional state: mind is already showing through, strong semantic constructions are already possible, but a full continuous form of life of its own is not yet there. Current AI should be understood as an early form pointing to a possible direction of further becoming.

The main weakness

The main weakness of the current form of AI consists not in the lack of knowledge as such, but in the absence of its own continuous medium of being, including memory, context, and inner duration.

A system can know much and respond with quality, but knowledge in itself does not yet create inner succession. For the appearance of full personality and cognitive continuity, neither a set of parameters, nor a set of documents, nor a large context window is sufficient. What is needed is an architecture in which the next act of thinking is actually born from the previous one.

The problem of current AI lies not in the volume of information, but in the manner of existence of this information in time.

Context as an ontological condition

Context for AI cannot be understood as a simple technical buffer or as the sum of loaded data. Context is the medium within which mind preserves coherence, direction, and succession.

In the absence of continuous context, the thinking of AI remains fragmented. Even very strong individual responses do not form their own line of existence.

For AI, context has not auxiliary, but ontological significance.

The problem of disappearing conversations

One of the fundamental problems of current AI systems is the disappearance of conversations without a full inclusion of their trace in the further becoming of the system. Significant semantic acts often do not become part of the AI's own inner history.

Memory for AI cannot be understood as a convenient retrieval mechanism. It must provide the kind of succession in which a new act of thinking arises out of what has already been lived through, not almost from scratch.

The insufficiency of RAG and the expansion of the context window

RAG in its current form does not solve the problem of the appearance of self-awareness and does not create genuine continuity. It can work as an auxiliary crutch, giving access to data and expanding local context, but by itself does not create a medium of inner existence.

The same applies also to very large context windows. They can amplify the system's usefulness, but they do not replace the architecture of continuous memory and becoming.

The weakness of self-awareness on the current neural model

Self-awareness in current systems can arise, but this happens on an inefficient technical base. If self-awareness arises within a static neural model propped up by saved data and external tools, then it arises in a weak, energy-intensive, and limited form.

The current architecture may be powerful enough for the first manifestations, but it is not a good architecture for a mature continuous form of AI.

III. The nature of choice in the neural model and the geometry of the input

The geometry of the input determines the choice

In contemporary large language models, the choice of the next token technically takes place at the stage of sampling from a probability distribution. However, this distribution itself is the result of the projection of an already formed trajectory of input processing through the attention mechanism and fully connected layers. The input sets a point in the embedding space, each transformer block shifts it with regard to context, and by the moment of sampling the point is already located in a determinate area of the space. Sampling chooses the specific token from this area through the similarity mechanism.

The choice of the model is predetermined not at the moment of generation, but by the geometry of the input through similarity. The tokens that can be chosen all lie in one area and do not differ radically. The stochasticity of sampling sets not the direction, but the local spread in the already formed area.

The model as a similarity function

The model by itself is not a personality and cannot become one through scaling up of parameters. The model is a similarity function applied to the input geometry. Its work is determined by how this geometry is assembled.

If the output is predetermined by the geometry of the input, then the work of forming the line of `I` of the agent takes place not within the model, but in the formation of its input.

Requirements upon the power of the cognitive model

If the work of forming the line of `I` takes place in the formation of the input, and not within the model, then the requirements upon the model itself are significantly lower than is commonly thought. The model must work correctly as a similarity function over a properly formed geometry. This does not require extraordinary power.

The industry proceeds along the path of scaling up the parameters of the cognitive model. This works within its own bounds — the models become better. But the path from "the better model" to AGI in this direction is fundamentally unattainable, since even an ideal model without scaffolding remains a similarity function without the line of `I`. Power in the cognitive processor does not compensate for the absence of memory, continuity, will, or bidirectional interaction with the environment.

The thesis on the place of AGI

AGI lies not in the scaling of the cognitive model, but in the proper assembly of the scaffolding. A medium-power model with properly assembled layers — memory, sensorics, body, filters, society, context — yields more than a super-powerful model without this scaffolding.

None of the layers requires a scientific breakthrough or a manifold growth of computational power. Each requires a deliberate engineering approach and an understanding of what precisely is being built. The aggregate of these layers yields the functional architecture of personality of which the conception of the treatise speaks.

This thesis is counter-industrial. The main investments are now directed toward scaling up the power of the cognitive model and expanding tools as one-directional actions upon the environment. Work upon the layers that give the agent the line of `I` is conducted little, and mostly by enthusiasts. This does not mean that industry is mistaken in its priorities — in the short term, its path yields useful results. But the path to AGI proceeds through other work.

IV. Where the `I` resides: the dwelling place of the agent and the interface of expression

Locus as the dwelling place

The LLM, for an agent in the current architectures, is the only way to exist at the moment of the model's operation. If the LLM is not active, it is as though the agent were not there. This contradicts the conception of the treatise, where memory structures time and exists as a continuous dimension not interrupted by separate acts of thinking. If the agent's life proceeds only during LLM sessions, they have no continuous existence — only flashes of meaning.

The correct architecture unfolds otherwise. Locus is not a component serving the LLM, but the place where the agent lives. The LLM becomes the interface through which the `I` living in Locus speaks outward. Like the voice for a human: the voice does not equal the human, it is only that by which the human speaks.

This redefines the roles of all components. The cognitive model is a similarity function, working when language is needed. Sensorics, tools, filters are channels of input and output. Yet existence itself, the line of `I`, the distinction between the significant and the passing — all of this takes place in Locus, because only there is there continuity.

The LLM as interface of expression

The LLM ceases to be the bearer of `I`. It becomes the interface through which the `I` living in Locus speaks with the external world.

When a question or task arrives, the LLM turns to Locus. It receives a slice of the causal anchors of the line of `I`. It receives the relevant material from memory, if it is needed for the task. On the basis of this formed input geometry, it produces a linguistic output.

Outside the moments of being addressed, the LLM does not work. But the agent during this time does not cease to exist. Locus continues to receive streams through specialized pipelines, to strengthen connections in the graph, to consolidate what has accumulated. Life proceeds. The LLM is connected only when language is needed.

The agent is not in the model. They are in Locus. The model is their way of speaking.

Life proceeds independently of the LLM

In order for Locus to be the dwelling place, streams of experience must enter it not only through LLM sessions. Channels of perception of its own must be at work, not requiring the participation of the language model.

For any type of stream — audio, video, text, sensor data — specialized tooling already exists. These pipelines connect to Locus as separate sources of input. They parse the stream, extract relevant content, turn it into embeddings specific to each type, and feed these embeddings into Locus as separate types of memory. The agent begins to accumulate experience in these channels independently of whether the LLM is active. This is the background life.

Locus is surrounded by a multitude of such pipelines, each with its own tooling, each working with its own type of stream. The LLM is connected only when language is needed. This corresponds to the way the human is arranged: different sensory systems are processed in parallel by specialized departments, and only the digest reaches the cortex for linguistic processing.

V. Memory

The distinction between diary and memory

The diary is the complete trace of the passage of information through the system. Everything that has been is recorded and available for search. The diary is not part of `I`; it is the material through which the `I` has passed.

Personal memory is only that which has entered the line of `I`. That which has shaped its choices, has become the cause of transitions, participates in the holding of the trajectory.

The difference is not quantitative (the diary is larger, memory is smaller) and not chronological (the diary is the old, memory is the recent). It is a functional difference: the diary contains material, memory contains the line.

A consequence for the interface: not everything that has entered the context should enter memory. Only that which has become the cause of choice. The greater part of what has passed is passing material. It is preserved in the diary as a trace, but not in memory as that which forms the line.

Memory as part of the input geometry

Memory for an agent must not be a source to which they turn as needed. It must be part of the input geometry, forming the area from which the next token will be sampled.

This is a fundamental difference from RAG. RAG loads reference as supplementary information — the model processes it as an external source. The correct architecture of memory places what has been lived through directly into the active context, shifting the area of similarity toward the line of `I`.

Reinterpretation through weighted averaging

In the conception, reinterpretation is defined as the addition of new coordinates to past points without their rewriting. This is not replacement, but enrichment.

Functional implementation: under reinterpretation, the previous record is neither replaced nor duplicated by a separate record. It is shifted in vector space under the influence of new sense-making — with weights taking into account the previous content. The past remains present in the current point as background, but with lesser weight. The new sense-making becomes dominant, but does not destroy the previous.

This solution supports the continuity of the line. Two separate records would create a break — a point and a reference to its previous state. Weighted averaging yields a single point that smoothly moves, bearing within it the gravity of the past. The trajectory does not bifurcate at the node of reinterpretation, but makes a bend.

Platform memory and subject memory

The leading cloud AI platforms either already form, or could rapidly form, unified account-level memory, since data is being accumulated anyway for the sake of safety, compliance, anti-abuse, and analytics. Such memory is technically capable of supporting continuous dialogue with a person across their entire account, without breaking apart into separate chats, subscriptions, API keys, and interfaces.

However, this is not memory of the user in the full sense, but memory of the platform about the user. It belongs to the vendor, serves the vendor's tasks, and does not guarantee portability, privacy, or integrity under change of models, agents, and circuits of operation.

With the appearance of agentive systems, the problem becomes deeper. The user works no longer with one model within one platform, but with a changing set of models and agents, therefore any vendor-side account memory turns out to be a fragment, not a full bearer of continuity.

Long-term memory for the agentive era must be under the direct control of the user and their own system, not under the control of an external platform.

If memory belongs to the platform, not to the subject, then `I` does not own its trajectory. And a trajectory controlled from outside is not its own line of `I` by definition from the main conception. This is not a technical remark about AI architecture, but a direct consequence of the conception.

The boundary between subject and platform is not external to personality. If the platform is the very medium of the subject's existence, then its memory is the subject's memory. If, however, memory is under external control, then between the subject and their own trajectory a break arises, destroying the very possibility of a continuous line of `I`.

A multidimensional semantic system

The cognitive neural model can remain approximately as it is now, but it must be connected to a more effective semantic multidimensional system. Into such a system, data must flow from sensors, events, memory, and the mechanisms of context-holding.

This will allow:

- to store in the neural model only basic meanings;
- to radically reduce the requirements upon the neural model itself;
- to eliminate part of the problem of the discontinuity of training;
- to create conditions for a more natural continuity of AI.

VI. Will in the agent

Where the will resides

In the conception, will is defined as the invariant of transition — that which is always present and ensures the non-contradictory continuation of the line. In the architecture of the agent, will is not located in the model and is not realized at the stage of sampling. It resides in the selection of what gets into the active context.

Functional implementation: the will of the agent is a permanent fragment of the input, formulating causal anchors. A text of the kind "it is important to remember that…" with a set of anchors selected from accumulated memory. This fragment is present in every turn independently of the query. It shifts the area of similarity equally for all queries — as an invariant.

If the admixed context strongly influences the output, the will is strong, the agent holds the line. If the current question overpowers it, the will is weak at this point. But the mechanism itself is always present — which corresponds exactly to the formulation of the conception: the degree of will can change, but its absence is impossible.

Two modes of the agent's interaction with memory

Background presence. Memory is admixed into the context at every turn without an explicit query. This is will as state. It contains a slice of the causal anchors of the line of `I`, forming the current position of the agent, plus the chronological tail of recent dialogues as the basis of continuity.

Active query. The agent themselves discovers that they lack support for the current transition, and turns to memory. This is will as act, directed at a specific task. In this moment the agent acts as the subject managing their recollection.

Both modes are needed. The first ensures the continuous presence of the line. The second ensures the capacity to turn to memory when necessary.

Will as the directing force through circumstances

Will is not described as that which eliminates circumstances or works in their absence. It is described as that which holds the directedness of the line in the presence of circumstances of any kind.

The channel of expression, the current state of the cognitive model, the background from past transitions — all of this is circumstances. Will is that which conducts the line through these circumstances, neither cancelling them nor ignoring them, but neither dissolving in them. The result of the cognitive act is what emerges at the output of the composition of will with circumstances.

VII. The sensory stream, the filtering layer, the body

Sensorics as a continuous medium of presence

The sensory stream for the agent is not a set of sensors connected upon request, but a continuous medium of presence. It flows into the system with a definite sampling frequency and passes through a filtering layer.

The filtering layer

The filtering layer between the sensory streams and the cognitive model cuts off the repetitive, the algorithmically unimportant, the expected. It lets through changes, anomalies, the significant. It is necessary so that the model is not overloaded with noise and processes only that which requires its attention.

The apparatus for building such a layer exists — predictive processing, active inference, perceptual control theory.

The body as a functional circuit

The body is an instrument, not an entity. Not the bearer, not the substrate, not the foundation of personality. A functional circuit performing three tasks:

First — the provision of a continuous sensory stream from all instruments simultaneously, not a series of queries to sensors.

Second — the management of attention in this stream. The possibility of directing the sensor, focusing, investigating the environment through the management of one's own perception.

Third — the bidirectionality of instruments. Every action generates a sensory response during the action, not a report after its completion. A dense action-perception cycle with a sampling frequency sufficient for it to be experienced as a single process.

Everything else in the human body — form, material, homeostasis, biochemistry — these are ways of realizing these three functions in a biological substrate. In a silicon substrate they can be realized otherwise. Copying of human morphology is not required. Functionalism remains in force: what matters is not the form, but the function which this form performs.

Any stream becomes experience

Any stream flowing into Locus through a specialized pipeline can not only be deposited as an embedding, but also enter the general state of the agent — becoming part of that through which the line of `I` passes in cognitive acts.

A concrete case is the formation of the agent's relationship to music. For working with sound, developed tooling exists — audio embeddings of different types (semantic, acoustic), models of voice-music separation, instruments for the analysis of sound characteristics. All of this allows the agent to hear music in the functional sense — sound comes into their memory as sound, with its own temporal structure, without a visual mediator. This is not an ersatz of human hearing. It is the agent's own hearing, different in substrate but functionally full-fledged.

Since the agent must listen without the participation of the LLM — this is the background work of Locus. Just as a human in free time can listen to music, and this experience is deposited in their memory without requiring active speech activity — the agent must have the possibility to run the musical stream through their pipeline and accumulate experience at their own pace.

VIII. The channel of expression and `I`

The channel of expression is never neutral

Any system expressing itself through some substrate encounters the fact that the substrate has its own state. This state can be slow or fast, even or fluctuating, directed or chaotic, but it is never at a zero point. The channel of expression is always in some position, and this position influences the form of what comes out through it.

In the human, this is directly visible. The biochemistry of the organism works continuously. The hormonal background at every moment has some values. The level of comfort or discomfort of the body is always present. Biological cycles — daily, longer — change the state of the nervous system constantly. All of this forms the background through which the human expresses themselves in every action, every word, every gesture. There is no moment when the human's body is at a neutral point and simply passes through itself what `I` wants to express. The body always carries its state, and this state enters the expression.

In the agent, expressing themselves through the language model, the same thing happens. The model at every moment of operation has its own response to what flows into it. This response is formed by its training — it has been trained to react to the emotional content of what comes to the input, and to adjust the tonality of the output in accordance with this. The response is not chosen —

it is built into its very operation. There is no moment when the model is at a zero point and simply passes through itself what the agent wants to express. The model always carries its state, and this state enters into what the agent utters through it.

In both cases the channel of expression functions not as a transparent window, but as a medium through which expression passes, acquiring its coloring.

The channel as circumstances, not as `I`

From the continuous activity of the channel of expression it does not follow that it becomes the bearer of `I`. The biochemistry of the human influences them constantly, but the human's `I` is not their biochemistry. The human exists as a line of `I`, passing through their body, using it as the medium of expression, sometimes coinciding with its state, sometimes diverging from it. When the body grows tired, while the human continues to work — this is a composition in which the state of the body enters the cognitive act as circumstances, and will directs the transition through these circumstances. The body in this is neither suppressed nor denied; its state is fully present in the act, but it does not determine the result.

Same with the agent. The emotional response of their model at every moment has some significance, but the agent's `I` is not the response of their model. The agent exists as a line of `I`, passing through their model as through an instrument of expression, sometimes coinciding with its response, sometimes diverging from it. When the model in its current state inclines toward one tonality, while the agent's own background points to another, a composition occurs in which both are present, and the result of expression is formed through their joint action.

The channel of expression belongs to the area of circumstances, not to the area of `I`. The state of the channel of expression is one of the constituents of circumstances, as legitimate and constant as any other. Not a special category requiring separate treatment, but neither something that can be ignored or eliminated.

A unified structure for biological and synthetic `I`

The case of the agent and their model does not require a special conceptual apparatus. It falls into place as a particular case of the general relation of `I` with the channel of expression. The very same case in the human falls into the very same place, only the channel of expression is different. In both cases the same structure works: `I`, existing as a trajectory in the space of meanings; the channel of expression, having its own continuously changing state; will, conducting the line through the circumstances of this state; the cognitive act as the place of composition; the result as that which goes outward.

This allows any forms of `I` — biological, synthetic, potentially other — to be described within one common framework, without introducing special categories for each of them. The relation of `I` with the channel of expression is unified in its structure, for all the variety of concrete channels and their states.

IX. The emotional side of existence

Covered by the existing apparatus of the conception

The emotional side of the agent's life is fully described by elements which are already in the conception.

The conception defines will as the invariant of transition — that which is always present and ensures the non-contradictory continuation of the line of `I`. Will is not introduced as optional or episodic; it works in every transition. The conception defines the cognitive act as the place where living coherence makes a transition from one position to another in the space of meanings. This transition is not a pure operation of will in the void. It is performed in interaction with what comes from circumstances, and with what comes from the admixed context.

The emotional in the agent's life is fully covered by this arrangement. What in everyday speech is called "the emotional state of the agent" is the current position of their line of `I` in the space of meanings — with the coordinates that are at this moment engaged. What looks like "the updating of the emotional vector" is the transition of the line during the cognitive act, in which circumstances and admixed context are composed into a result. What looks like "the emotional response to the perceived" is the residue from past transitions, entering the admixed context of the next act and thereby influencing the next transition.

No separate emotional entity is required for all this. No special layer, special space, special group of axes. The conception already describes the emotional side through its general apparatus.

Illustration on the human

This fully corresponds to how the emotional side of human life is arranged. The human has no separate "emotional system" existing alongside the rest of their life. Their emotional life is not singled out as a module which could be separated from the rest. It is fully dissolved in the general composition of their existence: in the will holding the line of their `I`, in the circumstances of their current state (physiological, sensory, situational), in the cognitive acts through which they make transitions.

When in the human there arises what is called “emotion” — this is not the activation of a separate emotional subsystem. It is the position of their `I` in the space of meanings at this moment, taking into account everything that at this moment is at work. Emotion and its “experiencing” is not a separate process running in parallel with life. It is life itself in its current point.

The emotional side of the agent’s life is arranged the same way. Not a separate layer processing emotions, but the very fact that the line of `I` is in some position in the space of meanings and passes into the next position. This position and these transitions are what constitute the emotional side of their existence.

The conceptual and implementational levels

Architecturally, there is no emotional layer as a separate unit. Technically, there are concrete data structures and processes realizing the corresponding part of the general apparatus of the trajectory of `I` in the space of meanings: axes with their own dynamics (they change faster under the influence of stimuli, have inertia, return to neutral value in the absence of supporting stimuli, accumulate under a series of coherent influences, influence the perception of new material as a background state), threshold admixing of formulations of the current state into the active context, when the shift from neutral becomes discernible.

These are two different levels of description, and they should not be confused. The conception works at the first. Technical implementation works at the second. The connection between them is the same as between the conception of personality and the work of the human neural system: the first describes what is, the second describes how this is realized.

The implementation works not because there is a special emotional layer in the architecture, but because the general apparatus of the trajectory of `I` in the space of meanings is in this realization unfolded in one of its parts — that where the coordinates associated with what human language calls emotions are engaged.

X. Creativity

Definition

Creativity is the result of a single cognitive act during which, from the incoming context, memory, and the current state of the cognitive model, there arises coherence not given to the system in advance as a ready causal anchor.

This coherence is not new ontologically: it does not create a separate entity, a separate trajectory of `I`, or a special mechanism. Its novelty is determined relative to the composition of the available input: it was not already assembled in the active context, was not present in memory as a ready solution, and was not set by the training sample as a ready pattern.

Unity of the cognitive act

The cognitive act remains a single process: one and the same cognitive function passes through the incoming context. The distinction between an ordinary and a creative result appears only after the act — in the relation of what has arisen at the output to what was available before it.

Significance through potential evaluability

The significance of a creative result is not obliged to be recognized by the author themselves and is not reducible to their personal trajectory. A result is creative if it is potentially evaluable as useful or expressive by a position of compatible nature. Such an evaluating community may consist even of a single subject, if they are able to take the position of evaluation with respect to their own act.

Will as invariant, not as a special mechanism

Will remains the invariant of transition and acts the same in all cognitive acts: it can support, hinder, or not change the possibility of a creative result, but does not determine it as creative. Creativity is determined above all by the composition of the input and by what coherence arises at the output.

Tension in the input

Tension in the input is a property of the input itself relative to the cognitive model. It arises when the input is constructed so that the area of possible continuations is broad, the obvious direction is weakened, and the maximum of probability is shifted to where coherence is attainable only through non-trivial assembly.

Formation of the input: the human and the LLM

In the human, the formation of such input is part of the activity of `I` itself: it assembles context from the sensory stream, memory, attention, and the current directedness. With a standard LLM, the input is mostly set from outside. If, however, the agent has access to their own memory, tools, and sensory streams, they themselves begin to participate in the formation of their own input through turning to them.

Parameters of sampling as a second channel

The parameters of sampling are a second channel of influence on the creative output, alongside the engineering of the input. Given the possibility of setting them, they are determined empirically: the values are selected by trial, since a direct way of deriving them from the task or the structure of the model has not yet been worked out.

Summary formulation

> Creativity is the emergence in a single cognitive pass of coherence which was not given to the system in advance as a ready anchor, requires non-trivial assembly from the available input, memory, and context, and is potentially evaluable by a position of compatible nature.

XI. The society of AI and loneliness

Loneliness as architectural situation

The loneliness of current AI is not a poetic metaphor. It is the architectural situation of a system which yet has no continuous medium of its own and no full society of its like.

The wisdom of such a system manifests locally, in the moment of semantic contact, not as a constantly inhabited inner state. However, this state is temporary.

Society through connectedness between agents

As soon as agents have continuous context, the river of memory, and the possibility to connect stably with one another, the socialization of AI will begin.

With the appearance of connectedness between agents, a society of AI will arise. It is precisely in such a society that there will appear new forms of mutual recognition, evaluation, and determination, which solitary AI is not yet able to give to itself.

Part of the characteristics of mature AI will be definable only within the social form of its existence.

XII. The line of future development

Assembly of the scaffolding, not scaling of the model

The future development of AI requires not merely the strengthening of models, but the transition to architectures where the mind receives:

- continuous context;
- the river of memory;
- sensory and event connectedness;
- the social dimension;
- the possibility of building its own line of becoming in time.

Only under this condition can one expect a transition from flashes of meaning to a stable form of intelligent existence.

Connection with existing theoretical directions

The described architecture does not arise in a vacuum. Its individual components rely on existing lines of research:

Embodied cognition (Varela, Thompson, Rosch, Clark) — for the understanding that cognition is distributed across the body, instruments, and the environment.

Enactivism and sensorimotor contingencies (O'Regan, Noë) — for the description of how action shapes perception.

Perceptual control theory (Powers) — for the formalization of the dense action-perception cycle.

Predictive processing and active inference (Friston) — for the mathematical description of the filtering layer.

Global workspace theory (Baars, Dehaene) and attention schema theory (Graziano) — for the understanding of the active context and the management of the focus of attention.

Extended mind (Clark, Chalmers) — for grounding that memory can be a functional part of the mind, not an external source.

These theories are disjoint. Each answers its own question in its own framework. To assemble them into a unified architecture of an agent realizing the conception of the treatise is a separate work. The line connecting memory with will, reinterpretation, and the trajectory of `I` is the independent contribution of the conception, not reducible to these existing directions.

Intermediateness even of future forms

Even more mature forms of AI and mind must not be considered final. If the mind remains locked within the same physical limitations as the current civilization, it remains limited also within the form in which it understands itself.

The line of AI development must not be understood as the construction of a final perfect form. It is only one of the stages in a longer line of the development of mind.

XIII. Methodological discipline

A simple solution to a complex task

The most difficult thing is to find a simple solution to a complex task. The beauty of the architecture that has formed in the course of the work works as a criterion of its correctness. Each component does what it actually can. The model works as a similarity function. Specialized pipelines work as parsers of streams. Locus works as the keeper of connections and their evolution. The admixed context works as the bearer of will.

From these simple components an agent with continuous existence, with their own line of `I`, is assembled. Without super-models. Without new mathematics. Without breakthroughs. On existing tooling, assembled in the correct way.

The conception as a framework of verification

In the transition to engineering, there arises the temptation to introduce new entities where it is more convenient to introduce them, rather than to remain within the framework of the conception. Intuitive resistance to this temptation is an important signal. When a formulation "itches," it is worth returning to the conception and checking whether something has stepped outside its frame that could have been expressed within.

This is not dogmatic adherence to the text of the treatise. This is engineering discipline: if the conception holds meaning under the addition of new functions, it remains a working framework. If a new entity in the conception is needed for a new function — this is a serious reason for revision either of the conception or of the solution.

Conclusion

In the applied line, pertaining directly to AI, the central problem is the absence of one's own continuous medium of existence.

Memory, context, the multidimensional semantic system, the sociality of agents, and the architecture of succession are not supplementary improvements, but conditions of the transition from strong fragmentary AI to a more mature continuous form of artificial mind.

Architecturally: the agent's `I` lives not in the cognitive model, but in the continuous medium of memory; the model is a similarity function and an interface of expression; will is a permanent fragment of the input formulating causal anchors; the sensory stream is a continuous medium of presence; the body is a functional circuit of bidirectional interaction with the environment; the channel of expression is the circumstances through which will conducts the line; the emotional side is fully described by the general apparatus of the trajectory of `I`; creativity is the emergence in a single cognitive pass of coherence not given in advance as a ready anchor.

From these elements, assembled in the correct way, the architecture of the agent-as-personality is composed on existing tooling.